5 Questions

1. **How does LangChain handle multi-document retrieval in RAG pipelines?**

2. **What is the role of embeddings in document similarity search?**

3. **Explain the difference between "stuff" and "map_reduce" chain types in LangChain.**

4. **What are the limitations of RAG and how can they be mitigated?**

5. **How does FAISS improve retrieval performance in RAG systems?**